

```
export CLUSTER_NAME="ml-gke"
```

```
gcloud container clusters create $CLUSTER_NAME \
  --enable-image-streaming \
  --addons=HttpLoadBalancing \
  --machine-type=e2-standard-2 \
  --shielded-secure-boot \
  --shielded-integrity-monitoring \
  --region=us-central1 \
  --num-nodes=1 \
  --enable-ip-alias \
  --release-channel=rapid \
  --node-locations=us-central1-a \
  --addons=RayOperator
```

```
kubectl apply -f - <<'EOF'
```

```
apiVersion: ray.io/v1
```

```
kind: RayJob
```

```
metadata:
```

```
  name: rayjob-sample
```

```
spec:
```

```
  # submissionMode specifies how RayJob submits the Ray job to the RayCluster.
```

```
  # The default value is "K8sJobMode", meaning RayJob will submit the Ray job via a
  submitter Kubernetes Job.
```

```
  # The alternative value is "HTTPMode", indicating that KubeRay will submit the Ray job by
  sending an HTTP request to the RayCluster.
```

```
  # submissionMode: "K8sJobMode"
```

```
  entrypoint: python /home/ray/samples/sample_code.py
```

```
  # shutdownAfterJobFinishes specifies whether the RayCluster should be deleted after the
  RayJob finishes. Default is false.
```

```
  # shutdownAfterJobFinishes: false
```

```
  # ttlSecondsAfterFinished specifies the number of seconds after which the RayCluster will
  be deleted after the RayJob finishes.
```

```
  # ttlSecondsAfterFinished: 10
```

```
  # activeDeadlineSeconds is the duration in seconds that the RayJob may be active before
```

```
  # KubeRay actively tries to terminate the RayJob; value must be positive integer.
```

```
  # activeDeadlineSeconds: 120
```

```
  # RuntimeEnvYAML represents the runtime environment configuration provided as a
  multi-line YAML string.
```

```
  # See https://docs.ray.io/en/latest/ray-core/handling-dependencies.html for details.
```

```
  # (New in KubeRay version 1.0.)
```

```
  runtimeEnvYAML: |
```

```
    pip:
```

```
      - requests==2.26.0
```

```

- pendulum==2.1.2
env_vars:
  counter_name: "test_counter"

# Suspend specifies whether the RayJob controller should create a RayCluster instance.
# If a job is applied with the suspend field set to true, the RayCluster will not be created and
we will wait for the transition to false.
# If the RayCluster is already created, it will be deleted. In the case of transition to false, a
new RayCluster will be created.
# suspend: false

# rayClusterSpec specifies the RayCluster instance to be created by the RayJob controller.
rayClusterSpec:
  rayVersion: '2.9.3' # should match the Ray version in the image of the containers
  # Ray head pod template
  headGroupSpec:
    # The `rayStartParams` are used to configure the `ray start` command.
    # See
https://github.com/ray-project/kuberay/blob/master/docs/guidance/rayStartParams.md for the
    default settings of `rayStartParams` in KubeRay.
    # See https://docs.ray.io/en/latest/cluster/cli.html#ray-start for all available options in
    `rayStartParams`.
    rayStartParams:
      dashboard-host: '0.0.0.0'
  #pod template
  template:
    spec:
      containers:
        - name: ray-head
          image: rayproject/ray:2.9.3
          ports:
            - containerPort: 6379
              name: gcs-server
            - containerPort: 8265 # Ray dashboard
              name: dashboard
            - containerPort: 10001
              name: client
      resources:
        limits:
          cpu: "1"
        requests:
          cpu: "200m"
        volumeMounts:
          - mountPath: /home/ray/samples
            name: code-sample
      volumes:
        # You set volumes at the Pod level, then mount them into containers inside that Pod
        - name: code-sample

```

```

configMap:
  # Provide the name of the ConfigMap you want to mount.
  name: ray-job-code-sample
  # An array of keys from the ConfigMap to create as files
  items:
    - key: sample_code.py
      path: sample_code.py
workerGroupSpecs:
  # the pod replicas in this group typed worker
  - replicas: 1
    minReplicas: 1
    maxReplicas: 5
    # logical group name, for this called small-group, also can be functional
    groupName: small-group
    # The `rayStartParams` are used to configure the `ray start` command.
    # See
    https://github.com/ray-project/kuberay/blob/master/docs/guidance/rayStartParams.md for the
    default settings of `rayStartParams` in KubeRay.
    # See https://docs.ray.io/en/latest/cluster/cli.html#ray-start for all available options in
    `rayStartParams`.
    rayStartParams: {}
    #pod template
    template:
      spec:
        containers:
          - name: ray-worker # must consist of lower case alphanumeric characters or '-', and
            must start and end with an alphanumeric character (e.g. 'my-name', or '123-abc')
            image: rayproject/ray:2.9.3
            lifecycle:
              preStop:
                exec:
                  command: [ "/bin/sh", "-c", "ray stop" ]
            resources:
              limits:
                cpu: "1"
              requests:
                cpu: "200m"
    # SubmitterPodTemplate is the template for the pod that will run the `ray job submit`
    command against the RayCluster.
    # If SubmitterPodTemplate is specified, the first container is assumed to be the submitter
    container.
    # submitterPodTemplate:
    #   spec:
    #     restartPolicy: Never
    #     containers:
    #       - name: my-custom-rayjob-submitter-pod
    #         image: rayproject/ray:2.9.3

```

```

#      # If Command is not specified, the correct command will be supplied at runtime
using the RayJob spec `entrypoint` field.
#      # Specifying Command is not recommended.
#      # command: ["sh", "-c", "ray job submit
--address=http://$RAY_DASHBOARD_ADDRESS
--submission-id=$RAY_JOB_SUBMISSION_ID -- echo hello world"]

```

```

#####Ray code sample#####

```

```

# this sample is from
https://docs.ray.io/en/latest/cluster/job-submission.html#quick-start-example
# it is mounted into the container and executed to show the Ray job at work

```

```

---
```

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: ray-job-code-sample
data:
  sample_code.py: |
    import ray
    import os
    import requests

    ray.init()

    @ray.remote
    class Counter:
        def __init__(self):
            # Used to verify runtimeEnv
            self.name = os.getenv("counter_name")
            assert self.name == "test_counter"
            self.counter = 0

        def inc(self):
            self.counter += 1

        def get_counter(self):
            return "{} got {}".format(self.name, self.counter)

    counter = Counter.remote()

    for _ in range(5):
        ray.get(counter.inc.remote())
        print(ray.get(counter.get_counter.remote()))

    # Verify that the correct runtime env was used for the job.
    assert requests.__version__ == "2.26.0"
EOF

```

Kubectl get svc

Kubectl port-forward svc/ 8265:8265

Kubectl get raycluster

Kubectl edit raycluster

kubectl apply -f - <<'EOF'

apiVersion: ray.io/v1

kind: RayService

metadata:

name: text-summarizer

spec:

serviceUnhealthySecondThreshold: 900 # Config for the health check threshold for Ray Serve applications. Default value is 900.

deploymentUnhealthySecondThreshold: 300 # Config for the health check threshold for Ray dashboard agent. Default value is 300.

serveConfigV2: |

applications:

- name: text_summarizer

import_path: text_summarizer.text_summarizer:deployment

runtime_env:

working_dir:

"https://github.com/ray-project/serve_config_examples/archive/refs/heads/master.zip"

rayClusterConfig:

rayVersion: '2.9.3' # Should match the Ray version in the image of the containers

#####headGroupSpecs#####

Ray head pod template.

headGroupSpec:

The `rayStartParams` are used to configure the `ray start` command.

See

<https://github.com/ray-project/kuberay/blob/master/docs/guidance/rayStartParams.md> for the default settings of `rayStartParams` in KubeRay.

See <https://docs.ray.io/en/latest/cluster/cli.html#ray-start> for all available options in `rayStartParams`.

rayStartParams:

dashboard-host: '0.0.0.0'

Pod template

template:

spec:

containers:

- name: ray-head

image: rayproject/ray-ml:2.9.3

ports:

```
- containerPort: 6379
  name: gcs
- containerPort: 8265
  name: dashboard
- containerPort: 10001
  name: client
- containerPort: 8000
  name: serve
volumeMounts:
- mountPath: /tmp/ray
  name: ray-logs
resources:
  limits:
    cpu: "2"
    memory: "8G"
  requests:
    cpu: "2"
    memory: "8G"
volumes:
- name: ray-logs
  emptyDir: {}
workerGroupSpecs:
# The pod replicas in this group typed worker
- replicas: 1
  minReplicas: 1
  maxReplicas: 10
  groupName: gpu-group
  rayStartParams: {}
# Pod template
template:
  spec:
    containers:
    - name: ray-worker
      image: rayproject/ray-ml:2.9.3-gpu
      resources:
        limits:
          cpu: 4
          memory: "16G"
          nvidia.com/gpu: 1
        requests:
          cpu: 3
          memory: "12G"
          nvidia.com/gpu: 1
    tolerations:
    - key: "nvidia.com/gpu"
      operator: "Exists"
      effect: "NoSchedule"
```

EOF

